

系統程式 HW3

410410060 資工二 林柔均

1. 請問 buffersize 分別是：0、-1、4KB、16KB、64KB、1MB、8MB的執行速度分別為何？

參數測試

Buffer_size: 0

```
A ~/De/h/Sy/hw4 > time ./fileperf input.txt output.txt 0 16:59:15
=====
CPU      99%
user     1.518 s
system   12.905 s
total    14.439 s
```

Buffer_size: -1

```
A ~/De/h/Sy/hw4 > time ./fileperf input.txt output.txt -1 14s 17:01:06
=====
CPU      97%
user     0.234 s
system   0.655 s
total    0.913 s
```

Buffer_size: 4KB = 4096 B

```
A ~/De/h/Sy/hw4 > time ./fileperf input.txt output.txt 4096 17:01:14
=====
CPU      90%
user     0.225 s
system   0.030 s
total    0.282 s
```

Buffer_size: 16KB = 16484 B

```
A ~/De/h/Sy/hw4 > time ./fileperf input.txt output.txt 16384 17:01:31
=====
CPU      88%
user     0.175 s
system   0.017 s
total    0.216 s
```

Buffer_size: 64KB = 65536 B

```
A ~/De/h/Sy/hw4 > time ./fileperf input.txt output.txt 65536 17:01:57
=====
CPU      77%
user     0.159 s
system   0.030 s
total    0.243 s
```

Buffer_size: 1MB = 1048576 B

```
A ~/De/h/Sy/hw4 > time ./fileperf input.txt output.txt 1048576 17:02:16
=====
CPU      88%
user     0.167 s
system   0.027 s
total    0.218 s
```

Buffer_size: 8MB = 8388608 B

```
A ~/De/h/Sy/hw4 > time ./fileperf input.txt output.txt 8388608 17:02:33
=====
CPU      80%
user     0.155 s
system   0.024 s
total    0.222 s
```

結論

Buffer_size 越大，程式速度越快。但當數值大到一定程度時加速效果變小，差異不大

2.使用 ltrace 觀察應用程式呼叫

我的 ArchLinux 裡沒有 ltrace 指令，要自行下載: `yay -S ltrace`

Buffer_size: 0 (unbuffered)

```
A ~/De/homework/SystemProgramming/hw4 > ltrace -c ./fileperf input.txt output.txt 0
% time      seconds  usecs/call   calls   function
-----
 82.53    29.642686      51    572463 fgetc
 17.31     6.219218      66    93804 fputs
  0.16     0.056003      66     846 fputc
  0.00     0.000466     233        2 fopen
  0.00     0.000304     152        2 setvbuf
  0.00     0.000205     205        1 atoll
  0.00     0.000180      90        2 fclose
-----
100.00    35.919062             667120 total
```

Buffer_size: -1 (linebuffered)

```
A ~/De/homework/SystemProgramming/hw4 > ltrace -c ./fileperf input.txt output.txt -1
% time      seconds    usecs/call   calls      function
-----
83.61      19.669146        34      572463  fgetc
16.03       3.770635        40      93804   fputs
0.36       0.083953        99       846    fputc
0.00       0.000291       145        2    fopen
0.00       0.000102        51         2    fclose
0.00       0.000080        40         2    setvbuf
0.00       0.000065        65         1    atoll
-----
100.00      23.524272                667120  total
```

Buffer_size: 4096

```
A ~/De/homework/SystemProgramming/hw4 > ltrace -c ./fileperf input.txt output.txt 4096
% time      seconds    usecs/call   calls      function
-----
85.68      19.008860        33      572463  fgetc
14.18       3.146580        33      93804   fputs
0.13       0.028371        33       846    fputc
0.00       0.000533       266         2    fopen
0.00       0.000132        66         2    setvbuf
0.00       0.000126        63         2    fclose
0.00       0.000092        92         1    atoll
-----
100.00      22.184694                667120  total
```

3. 使用 strace 觀察你的應用程式呼叫「作業系統的情況」

strace 也要自行下載: `yay -S strace`

Buffer_size: 0 (unbuffered)

```
A ~/Desktop/homework/SystemProgramming/hw4 > strace -c ./fileperf input.txt output.txt 0
% time      seconds    usecs/call   calls      errors  syscall
-----
66.73      1.146158         2      572464          read
33.26      0.571308         6      94650          write
0.01      0.000123        30         4          close
0.00      0.000000         0         8          mmap
0.00      0.000000         0         3          mprotect
0.00      0.000000         0         1          munmap
0.00      0.000000         0         3          brk
0.00      0.000000         0         2          pread64
0.00      0.000000         0         1          1 access
0.00      0.000000         0         1          execve
0.00      0.000000         0         2          1 arch_prctl
0.00      0.000000         0         1          set_tid_address
0.00      0.000000         0         4          openat
0.00      0.000000         0         2          newfstatat
0.00      0.000000         0         1          set_robust_list
0.00      0.000000         0         1          prlimit64
0.00      0.000000         0         1          getrandom
0.00      0.000000         0         1          rseq
-----
100.00      1.717589         2      667150          2 total
```

Buffer_size: -1 (linebuffered)

```
A ~/De/homework/SystemProgramming/hw4 > strace -c ./fileperf input.txt output.txt -1
```

% time	seconds	usecs/call	calls	errors	syscall
98.68	0.052162	4	10692		write
1.01	0.000533	3	142		read
0.32	0.000167	41	4		close
0.00	0.000000	0	8		mmap
0.00	0.000000	0	3		mprotect
0.00	0.000000	0	1		munmap
0.00	0.000000	0	3		brk
0.00	0.000000	0	2		pread64
0.00	0.000000	0	1	1	access
0.00	0.000000	0	1		execve
0.00	0.000000	0	2	1	arch_prctl
0.00	0.000000	0	1		set_tid_address
0.00	0.000000	0	4		openat
0.00	0.000000	0	4		newfstatat
0.00	0.000000	0	1		set_robust_list
0.00	0.000000	0	1		prlimit64
0.00	0.000000	0	1		getrandom
0.00	0.000000	0	1		rseq
100.00	0.052862	4	10872	2	total

Buffer_size: 4096

```
A ~/Desktop/homework/SystemProgramming/hw4 > strace -c ./fileperf input.txt output.txt 4096
```

% time	seconds	usecs/call	calls	errors	syscall
0.00	0.000000	0	142		read
0.00	0.000000	0	140		write
0.00	0.000000	0	4		close
0.00	0.000000	0	8		mmap
0.00	0.000000	0	3		mprotect
0.00	0.000000	0	1		munmap
0.00	0.000000	0	3		brk
0.00	0.000000	0	2		pread64
0.00	0.000000	0	1	1	access
0.00	0.000000	0	1		execve
0.00	0.000000	0	2	1	arch_prctl
0.00	0.000000	0	1		set_tid_address
0.00	0.000000	0	4		openat
0.00	0.000000	0	4		newfstatat
0.00	0.000000	0	1		set_robust_list
0.00	0.000000	0	1		prlimit64
0.00	0.000000	0	1		getrandom
0.00	0.000000	0	1		rseq
100.00	0.000000	0	320	2	total

ltrace & strace

ltrace & strace 的說明(Google 找到的):

ltrace的功能是能够跟踪进程的库函数调用，它会显现出哪个库函数被调用，而strace则是跟踪程序的每个系统调用。ltrace与strace使用的技术大体相同，但ltrace在对支持fork和clone方面，不如strace。strace在收到fork和clone等系统调用后，做了相应的处理，而ltrace没有。

版权声明：本文为CSDN博主「运维猫（运维开发）」的原创文章，遵循CC 4.0 BY-SA版权协议，转载请附上原文出处链接及本声明。

原文链接：<https://blog.csdn.net/yunweimao/article/details/106688073>

4. 有辦法根據 2 和 3 分析一下「呼叫作業系統核心函數（system call）」和「函數庫呼叫」的「成本」差異嗎？

整理數據:

- ltrace: 三個的 fgetc 的呼叫次數都是 572463, fputs 都是 93804 次

buffer size	0	-1	4096
time(second)	35	23	22

- strace:

buffer size	0	-1	4096
call times(read function)	572464	10692	142
time(second)	1.71	0.05	0.0

- 結論

system call 的時間和呼叫次數 都會隨 buffer size 而改變，function call 的呼叫次數並不影響

而呼叫和時間的成本:

time / call times	0	-1	4096
function call	$35/572464 = 0.00006$	$23/572464$	$22/572464$
system call	$1.71/572464 = 0.000002$	$0.05/10692$	$0.0/142$

可見 system call 的成本較低

註: 有參考學長姐的作業